

128-token prompts; 64 new tokens; unmerged adapters

FP32 compute

FP16 autocast

Generated tokens / second (batch aggregate)

